

The BFSS Model: Theory and Simulation

Zakaria Daoudi

September 27, 2026

Abstract

The BFSS model, proposed by Banks, Fischler, Shenker, and Susskind, is a conjectured non-perturbative formulation of M-theory in terms of the quantum mechanics of $N \times N$ matrices describing D0-branes. These notes give a pedagogical introduction to the model. We first review its theoretical underpinnings: its origin in the dimensional reduction of ten-dimensional super Yang–Mills theory, its symmetries, and its thermodynamics and gauge/gravity interpretation. We then explain in detail how the bosonic part of the model is simulated at finite temperature with the Hybrid Monte Carlo algorithm on a Euclidean time lattice, including the static diagonal gauge, the Faddeev–Popov term, the leapfrog integrator, the Metropolis test, and the measurement of observables such as the Polyakov loop and the internal energy. The discussion is illustrated with a short, self-contained Fortran code and a small temperature scan exhibiting the confinement/deconfinement transition. The notes are aimed at graduate students with a first course in string theory.

Contents

1	Introduction	2
2	Theoretical Background	2
2.1	M-Theory and Matrix Models	2
2.2	The BFSS Hamiltonian	3
2.3	Symmetries and Gauge Invariance	4
2.4	Finite Temperature and the Path Integral	5
3	Lattice Formulation of the Bosonic Model	5
4	Numerical Simulation: Hybrid Monte Carlo	7
4.1	The Fictitious Hamiltonian	7
4.2	Momentum Generation	7
4.3	Leapfrog Integration	8
4.4	Metropolis Test	9
4.5	Tuning and Checks	9

5	Observables	10
5.1	The Polyakov Loop	10
5.2	Energy	10
5.3	Extent of Space	10
5.4	Statistical Errors	11
6	Numerical Results	11
6.1	Checks of the Algorithm	11
6.2	Temperature Scan	11
7	Discussion and Outlook	12

1 Introduction

The BFSS model emerged in the mid-1990s as a candidate for a non-perturbative formulation of M-theory [1]. Based on the idea of describing D0-branes, point-like objects in type IIA string theory, by large $N \times N$ matrices, the model reinterprets M-theory as a quantum mechanical system. It provides a bridge between string theory, supersymmetric gauge theories, and quantum gravity, and it is one of the simplest settings in which gauge/gravity duality can be tested directly by numerical simulation [2–6].

These notes are organized as follows. In Section 2 we review the theoretical background of the BFSS model. Section 3 formulates the bosonic part of the model at finite temperature on a Euclidean time lattice. Section 4 describes the Hybrid Monte Carlo algorithm used to simulate it, and Section 5 discusses the physical observables, notably the Polyakov loop and the internal energy. Section 6 presents a small numerical study with the accompanying code, and Section 7 contains a discussion and an outlook.

2 Theoretical Background

2.1 M-Theory and Matrix Models

M-theory emerged in the mid-1990s as the framework unifying the five consistent superstring theories. It is an eleven-dimensional theory whose low-energy limit is eleven-dimensional supergravity and which, upon compactification, reduces to the known string theories. In particular, type IIA string theory is M-theory compactified on a circle of radius $R_{11} = g_s \ell_s$, where g_s is the string coupling and ℓ_s the string length. M-theory contains extended objects, the membrane (M2-brane) and the five-brane (M5-brane), and its Kaluza–Klein modes along the eleventh dimension, carrying momentum $p_{11} = n/R_{11}$, are identified with bound states of n D0-branes of type IIA string theory, whose mass is $1/(g_s \ell_s)$ [7].

The low-energy dynamics of N D0-branes is described by the dimensional reduction of ten-dimensional $\mathcal{N} = 1$ super Yang–Mills theory with gauge group $U(N)$ down to $0 + 1$ dimensions. The spatial components of the gauge field become nine Hermitian $N \times N$ matrices X_i ($i = 1, \dots, 9$), while the temporal component A becomes a non-dynamical gauge field that acts as a Lagrange multiplier enforcing the Gauss law constraint. When the matrices

commute they can be diagonalized simultaneously, and their eigenvalues are the positions of the N D0-branes in the nine transverse dimensions. The off-diagonal elements describe open strings stretched between different D0-branes; when they are excited, the matrices no longer commute and the notion of a definite position is lost. This non-commutativity is at the heart of the idea that classical geometry can emerge from an underlying quantum mechanical system.

The same matrix quantum mechanics had appeared earlier as the regularized description of the eleven-dimensional supermembrane [8]. Banks, Fischler, Shenker, and Susskind conjectured that, in the limit $N \rightarrow \infty$, this quantum mechanics provides a complete description of M-theory in the infinite momentum frame, with N units of momentum $p_{11} = N/R_{11}$ [1]. Susskind later proposed that at finite N the model describes the discrete light-cone quantization of M-theory [9]. For reviews, see e.g. [10].

An important motivation behind the BFSS proposal is its potential to describe the microscopic structure of black holes. At finite temperature and strong coupling, the model is dual to a black hole made of D0-branes in type IIA supergravity [2], and its thermodynamics reproduces the Bekenstein–Hawking entropy of that black hole. This has been verified by Monte Carlo simulations of the supersymmetric model [3–6], making the BFSS model a laboratory for studying the quantum properties of black holes from first principles.

2.2 The BFSS Hamiltonian

In units where the Yang–Mills coupling is absorbed into the fields, the bosonic part of the BFSS Hamiltonian is

$$H_B = \frac{1}{2} \text{Tr} (P_i P_i) - \frac{1}{4} \text{Tr} ([X_i, X_j][X_i, X_j]), \quad (1)$$

where repeated indices are summed over $i, j = 1, \dots, 9$, and P_i are the momenta conjugate to X_i . The commutator

$$[X_i, X_j] = X_i X_j - X_j X_i$$

measures the non-commutativity of different directions. Because $[X_i, X_j]$ is anti-Hermitian, $-\text{Tr}[X_i, X_j]^2 \geq 0$, so the potential energy is non-negative. The first term in Eq. (1) is the kinetic energy of the D0-branes and of the strings between them, while the second is the potential energy arising from their non-Abelian interactions.

The quartic potential vanishes exactly when all the matrices commute. The classical theory therefore has *flat directions*: N D0-branes placed at arbitrary positions, $X_i = \text{diag}(x_i^{(1)}, \dots, x_i^{(N)})$, cost no energy. Expanding around such a configuration, the off-diagonal element $(X_i)_{ab}$ acquires a mass $|\vec{x}^{(a)} - \vec{x}^{(b)}|$, which is the energy of a string stretched between D0-branes a and b . What happens to these flat directions quantum mechanically depends crucially on supersymmetry. In the purely bosonic model, the zero-point energy of the stretched strings generates a potential that grows with the separation, lifting the flat directions and binding the D0-branes together. In the supersymmetric model, the bosonic and fermionic zero-point energies cancel, the flat directions survive, and the spectrum is continuous. For D0-branes moving with relative velocity v at separation r , the leading one-loop effective potential is proportional to v^4/r^7 , which reproduces graviton scattering in eleven-dimensional supergravity. This was one of the first nontrivial checks of the conjecture [1, 10].

It is instructive to see how the coupling enters. Restoring the 't Hooft coupling $\lambda = g_{\text{YM}}^2 N$, the Euclidean action at temperature $T = 1/\beta$ reads

$$S_E = \frac{N}{\lambda} \int_0^\beta dt \operatorname{Tr} \left(\frac{1}{2} (D_t X_i)^2 - \frac{1}{4} [X_i, X_j]^2 \right) + S_F, \quad D_t X_i = \partial_t X_i - i[A, X_i], \quad (2)$$

where S_F contains the fermions. In $0+1$ dimensions, λ has dimensions of $(\text{energy})^3$. Under the rescaling $t \rightarrow t/s$, $X_i \rightarrow sX_i$, $A \rightarrow sA$, the bosonic action transforms as $S_E \rightarrow s^3 S_E$, which can be compensated by $\lambda \rightarrow s^3 \lambda$. Physical quantities therefore depend only on the dimensionless effective coupling

$$\lambda_{\text{eff}} = \frac{\lambda}{T^3}. \quad (3)$$

The model is weakly coupled at high temperature and strongly coupled at low temperature, where gravity is expected to emerge. In the rest of these notes we set $\lambda = 1$, so that temperatures are measured in units of $\lambda^{1/3}$.

The full BFSS Hamiltonian also includes fermionic terms, which are essential for supersymmetry:

$$H_F = \frac{1}{2} \operatorname{Tr} (\psi^T \gamma_i [X_i, \psi]), \quad (4)$$

where ψ is a sixteen-component real (Majorana) spinor of $SO(9)$ whose components are Hermitian $N \times N$ matrices, and γ_i are the real symmetric 16×16 gamma matrices of $SO(9)$ (the overall normalization depends on conventions). The model has sixteen supercharges. The supersymmetric cancellation between bosonic and fermionic contributions is what keeps the flat directions flat and what makes the gauge/gravity interpretation possible. Many numerical studies, including the one presented here, focus on the bosonic sector, which is simpler to simulate yet already displays non-trivial phase structure.

2.3 Symmetries and Gauge Invariance

The Hamiltonian (1) is invariant under time-dependent gauge transformations

$$X_i(t) \rightarrow U(t) X_i(t) U^\dagger(t), \quad A(t) \rightarrow U(t) A(t) U^\dagger(t) + i U(t) \partial_t U^\dagger(t), \quad (5)$$

where $U(t) \in U(N)$. The permutations of the D0-branes are a subgroup of $U(N)$, which reflects the fact that D0-branes are identical particles. Physical quantities must be gauge invariant, and are therefore built from traces of products of matrices, or from the eigenvalue distributions of the X_i .

Varying the action with respect to A gives the Gauss law constraint. In terms of the components $X_i = X_i^a T^a$ along the generators T^a of $U(N)$ with structure constants f^{abc} , it reads

$$G^a \equiv f^{abc} \left(X_i^b P_i^c - \frac{i}{2} \psi_\alpha^b \psi_\alpha^c \right), \quad G^a | \text{phys} \rangle = 0. \quad (6)$$

Physical states are therefore gauge singlets. At finite temperature this constraint is implemented by the integral over A in the path integral. The gauge freedom cannot be used to set $A = 0$ entirely: with periodic boundary conditions in Euclidean time, the gauge-invariant

holonomy $\mathcal{P} \exp(i \int_0^\beta A dt)$ remains. As we shall see in Section 3, the most one can do is to make A constant and diagonal, at the cost of a Faddeev–Popov determinant.

Beyond gauge invariance, the model has a global $SO(9)$ symmetry that rotates the nine matrices,

$$X_i \rightarrow R_{ij} X_j, \quad R \in SO(9), \quad (7)$$

and which descends from the ten-dimensional Lorentz symmetry. The fermions transform as spinors of $SO(9)$. Finally, the model has sixteen supersymmetries, which mix X_i and ψ and lead to the cancellations of quantum corrections discussed above.

The interplay between these symmetries has direct consequences for simulations. Monte Carlo methods require gauge fixing, or at least gauge-invariant observables. In addition, the bosonic model is not an approximation to the supersymmetric one in any controlled sense; it is a different theory with its own phase diagram. Its thermal behaviour is nonetheless interesting in its own right, and it serves as the natural first step towards simulating the full model.

2.4 Finite Temperature and the Path Integral

The thermodynamics of the model follows from the Euclidean partition function

$$Z(\beta) = \int \mathcal{D}A \mathcal{D}X \mathcal{D}\psi e^{-S_E[A, X, \psi]}, \quad (8)$$

where the bosons are periodic and the fermions antiperiodic in Euclidean time, with period $\beta = 1/T$. Integrating over A projects onto gauge-invariant states.

In the supersymmetric theory, the strongly coupled regime $T \ll \lambda^{1/3}$ is described by a black hole in type IIA supergravity made of N D0-branes [2]. Its thermodynamics predicts

$$\frac{E}{N^2} = 7.41 T^{14/5} \quad (\lambda = 1), \quad (9)$$

and this prediction, including its leading α' corrections, has been confirmed by Monte Carlo simulations [3–6].

In the bosonic model, the relevant physics is instead a confinement/deconfinement transition, familiar from gauge theories. Its order parameter is the Polyakov loop, defined in Section 5. At low temperature, the eigenvalues of the holonomy are distributed uniformly around the circle and the Polyakov loop vanishes at large N (confined phase). At high temperature they clump together and the Polyakov loop is non-zero (deconfined phase). Kawahara, Nishimura, and Takeuchi found the transition at $T \approx 0.9$ [11], and more recent large- N studies indicate a single first-order transition at $T_c \approx 0.885$ – 0.89 [12]. In the confined phase, observables are essentially independent of the temperature at large N . For example, $E/N^2 \simeq 6.70$ and $R^2 \simeq 2.29$ at $N = 32$ [11] (see Section 5 for the definitions). These numbers provide a useful benchmark for the code described below.

3 Lattice Formulation of the Bosonic Model

We discretize the Euclidean time circle into L sites, $t = 1, \dots, L$, with lattice spacing $a = \beta/L$ and periodic boundary conditions $X_i(L+1) = X_i(1)$. We keep a general number D of

matrices ($D = 9$ for the BFSS model) and set $\lambda = 1$. The trace part of each X_i describes the centre of mass of the D0-branes. It does not enter the commutators, so it decouples as a free particle, and it would be an exactly flat direction of the simulation. We therefore take the X_i traceless, i.e. we work with $SU(N)$ rather than $U(N)$.

Static diagonal gauge. On the lattice, the gauge field lives on the links as unitary matrices $U(t)$, and the covariant derivative becomes

$$D_t X_i \rightarrow \frac{1}{a} (U(t) X_i(t+1) U^\dagger(t) - X_i(t)). \quad (10)$$

Using gauge transformations, all the links but one can be set to the identity, and the remaining one, the holonomy, can be diagonalized. It is convenient to spread the holonomy evenly over all links, which gives the *static diagonal gauge*

$$U(t) = U = \text{diag}(e^{i\alpha_1/L}, \dots, e^{i\alpha_N/L}), \quad \text{i.e.} \quad A = \frac{1}{\beta} \text{diag}(\alpha_1, \dots, \alpha_N). \quad (11)$$

The change of variables from the $U(N)$ Haar measure of the holonomy to its eigenvalues produces the Vandermonde-like Faddeev–Popov determinant $\prod_{k < l} \sin^2(\frac{\alpha_k - \alpha_l}{2})$. We include it in the action as

$$S_{\text{FP}} = - \sum_{k < l} \ln \sin^2 \left(\frac{\alpha_k - \alpha_l}{2} \right). \quad (12)$$

This term repels the angles from each other and is responsible for the uniform distribution of the α_k in the confined phase.

The lattice action. Writing $\text{Tr}(UXU^\dagger Y) = \sum_{k,l} e^{i(\alpha_k - \alpha_l)/L} X_{kl} Y_{lk}$, the gauge-fixed lattice action is

$$\begin{aligned} S = & \frac{N}{a} \sum_{t=1}^L \sum_{i=1}^D [\text{Tr} X_i(t)^2 - \text{Tr}(U X_i(t+1) U^\dagger X_i(t))] \\ & - \frac{Na}{4} \sum_{t=1}^L \sum_{i,j=1}^D \text{Tr}[X_i(t), X_j(t)]^2 + S_{\text{FP}}[\alpha], \end{aligned} \quad (13)$$

which reproduces Eq. (2) (without fermions) as $a \rightarrow 0$. A mass term $\frac{Nam^2}{2} \sum_{t,i} \text{Tr} X_i(t)^2$ can be added if desired; we set $m = 0$. The same discretization is used in Refs. [11, 13].

Range of the angles. The angles α_k are defined modulo 2π : shifting a single α_k by 2π is a large gauge transformation, under which the matrices pick up periodic phases $X_{kl}(t) \rightarrow e^{\mp 2\pi i t/L} X_{kl}(t)$ that leave the action and all gauge-invariant observables unchanged. Ref. [11] restricts each angle to $-\pi < \alpha_k \leq \pi$. Since the action depends only on the differences $\alpha_k - \alpha_l$, one can equivalently let the overall translation of the angles run freely and require

$$\max_k \alpha_k - \min_k \alpha_k < 2\pi, \quad (14)$$

which is what we do. Both prescriptions give each configuration of eigenvalues on the circle the same weight and therefore the same expectation values. In the simulation, a proposed configuration that violates Eq. (14) is simply rejected.

4 Numerical Simulation: Hybrid Monte Carlo

Our goal is to generate configurations (X, α) distributed according to $e^{-S[X, \alpha]}$, with S given by Eq. (13), and to estimate observables as averages over these configurations. With DN^2L real degrees of freedom coupled through the commutator term, local updates are inefficient. The Hybrid Monte Carlo (HMC) algorithm [14, 15] instead updates all variables at once, by combining deterministic molecular dynamics in a fictitious time with a Metropolis test.

4.1 The Fictitious Hamiltonian

We introduce a Hermitian momentum matrix $P_i(t)$ conjugate to each $X_i(t)$ and a real momentum p_k conjugate to each α_k , and define the fictitious Hamiltonian

$$\mathcal{H} = \frac{1}{2} \sum_{t,i} \text{Tr } P_i(t)^2 + \frac{1}{2} \sum_k p_k^2 + S[X, \alpha]. \quad (15)$$

This is *not* the physical Hamiltonian (1). The momenta are auxiliary variables, and since $e^{-\mathcal{H}}$ factorizes into a Gaussian in the momenta times e^{-S} , sampling from $e^{-\mathcal{H}}$ and discarding the momenta produces the desired distribution. Hamilton's equations in the fictitious time τ are

$$\frac{dX_i}{d\tau} = P_i, \quad \frac{dP_i}{d\tau} = -F_i, \quad \frac{d\alpha_k}{d\tau} = p_k, \quad \frac{dp_k}{d\tau} = -f_k, \quad (16)$$

with forces $(F_i)_{kl} \equiv \partial S / \partial (X_i)_{lk}$ and $f_k \equiv \partial S / \partial \alpha_k$. From Eq. (13) we find

$$F_i(t) = \frac{N}{a} \left[2X_i(t) - UX_i(t+1)U^\dagger - U^\dagger X_i(t-1)U \right] - Na \sum_j [X_j(t), [X_i(t), X_j(t)]], \quad (17)$$

$$f_m = \frac{N}{aL} \text{Im} \sum_l \left[e^{i(\alpha_m - \alpha_l)/L} C_{ml} - e^{i(\alpha_l - \alpha_m)/L} C_{lm} \right] - \sum_{l \neq m} \cot \left(\frac{\alpha_m - \alpha_l}{2} \right), \quad (18)$$

where $C_{kl} = \sum_{t,i} X_i(t+1)_{kl} X_i(t)_{lk}$. The force $F_i(t)$ is Hermitian, so the evolution keeps X_i Hermitian. In practice, analytic forces should always be checked against finite differences of the action. The accompanying code does this when run with `test_force=.true..`

4.2 Momentum Generation

At the start of each trajectory, the momenta are drawn from the Gaussian distribution $e^{-\frac{1}{2} \text{Tr } P^2 - \frac{1}{2} \sum_k p_k^2}$. Since $\text{Tr } P^2 = \sum_k P_{kk}^2 + 2 \sum_{k < l} |P_{kl}|^2$, the diagonal elements of P are real Gaussian numbers with unit variance, while the real and imaginary parts of each off-diagonal element P_{kl} ($k < l$) have variance 1/2, with $P_{lk} = P_{kl}^*$. Finally, the trace is subtracted, $P \rightarrow P - \frac{1}{N}(\text{Tr } P)\mathbb{1}$. This is an orthogonal projection with respect to the metric $\text{Tr } P^2$, so it produces exactly the Gaussian distribution on traceless matrices. Since the forces (17) are traceless when X is, the matrices remain traceless throughout the evolution.

Gaussian numbers are generated from uniform ones with the Box–Muller transform [17]. Two independent uniform random numbers $u_1, u_2 \in (0, 1)$ are mapped to two independent standard normal numbers

$$g_1 = \sqrt{-2 \ln u_1} \cos(2\pi u_2), \quad g_2 = \sqrt{-2 \ln u_1} \sin(2\pi u_2). \quad (19)$$

The uniform numbers are produced by the `ran2` generator of Ref. [18], which combines two linear congruential generators with a shuffle and has a period of about 2×10^{18} . It must be initialized with a negative seed. In Fortran:

```
subroutine gauss_pair(g1, g2)
  real(dp), intent(out) :: g1, g2
  real(dp) :: r, ph
  r = sqrt(-2.0_dp*log(1.0_dp - ran2(idum)))
  ph = 2.0_dp*pi*ran2(idum)
  g1 = r*cos(ph)
  g2 = r*sin(ph)
end subroutine gauss_pair
```

(using $1 - u$ instead of u avoids $\ln 0$).

4.3 Leapfrog Integration

Hamilton's equations (16) are integrated numerically with the leapfrog integrator. One step of size $\Delta\tau$ consists of a half step for the fields, a full step for the momenta with the force evaluated at the midpoint, and another half step for the fields:

$$X_i\left(\tau + \frac{\Delta\tau}{2}\right) = X_i(\tau) + \frac{\Delta\tau}{2}P_i(\tau), \quad (20)$$

$$P_i(\tau + \Delta\tau) = P_i(\tau) - \Delta\tau F_i\left(\tau + \frac{\Delta\tau}{2}\right), \quad (21)$$

$$X_i(\tau + \Delta\tau) = X_i\left(\tau + \frac{\Delta\tau}{2}\right) + \frac{\Delta\tau}{2}P_i(\tau + \Delta\tau), \quad (22)$$

and similarly for (α_k, p_k) . A trajectory consists of n_τ such steps, and consecutive half steps for the fields are merged into full steps, so that each step costs one force evaluation:

```
X = X + 0.5*dtau*P           ! first half step
do step = 1, ntau
  call force(X, alpha, FX, Fa)
  P = P - dtau*FX           ! full step for the momenta
  if (step < ntau) X = X + dtau*P   ! two merged half steps
end do
X = X + 0.5*dtau*P           ! last half step
```

(with the same updates for `alpha` and its momenta). The leapfrog integrator has two properties that are essential for HMC. It is *time reversible*: integrating forward, flipping the sign of the momenta, and integrating forward again returns to the initial configuration. It is also *symplectic*, and in particular preserves the phase-space volume. It does not conserve $\mathcal{H}$ exactly. However, it exactly conserves a nearby “shadow” Hamiltonian that differs from $\mathcal{H}$ by $O(\Delta\tau^2)$ [19]. The energy violation $\Delta\mathcal{H}$ at the end of a trajectory of fixed length $n_\tau\Delta\tau$ therefore scales as $\Delta\tau^2$ and does not grow secularly with the trajectory length.

4.4 Metropolis Test

The integration errors are corrected exactly by a Metropolis test [16]. After a trajectory, the change in the fictitious Hamiltonian,

$$\Delta\mathcal{H} = \mathcal{H}_{\text{final}} - \mathcal{H}_{\text{initial}},$$

is computed, and the new configuration is accepted with probability

$$P_{\text{accept}} = \min(1, e^{-\Delta\mathcal{H}}). \quad (23)$$

If the new configuration is rejected, the old one is kept (and counted again). Together with reversibility and volume preservation, this ensures detailed balance with respect to $e^{-\mathcal{H}}$, so the algorithm is exact: the only effect of a large $\Delta\tau$ is a lower acceptance rate. In our case, a configuration that violates the constraint (14) is always rejected. In code:

```
dH = hamiltonian(X, alpha, P, Pa) - H_init
if (dH <= 0) then
    accepted = .true.
else
    accepted = (ran2(idum) < exp(-dH))
end if
if (.not. accepted) then
    X = X_old; alpha = alpha_old
end if
```

4.5 Tuning and Checks

The combination of the steps above (Gaussian momenta, a leapfrog trajectory, and a Metropolis test) is the HMC algorithm [14]. It proposes global updates of all DN^2L variables while keeping the acceptance rate high. Two parameters must be tuned: the step size $\Delta\tau$, and the trajectory length $n_\tau\Delta\tau$. A longer trajectory reduces autocorrelations but costs more. A common rule of thumb is to aim for an acceptance rate of about 70–90%. The largest force scale is set by the kinetic term, whose highest frequency is of order $2\sqrt{N/a}$, so $\Delta\tau$ must shrink as N and L grow. The step size must be fixed for the whole run, or at least not depend on the current configuration: adapting it along the Markov chain on the fly would spoil detailed balance. It is therefore tuned during thermalization.

Several exact identities provide stringent checks of an implementation:

- the forces agree with numerical derivatives of the action;
- integrating forward and then backward recovers the initial configuration to machine precision;
- at fixed trajectory length, $|\Delta\mathcal{H}| \propto \Delta\tau^2$;
- the equilibrium average $\langle e^{-\Delta\mathcal{H}} \rangle$ equals 1, which follows from volume preservation and reversibility alone.

Different step sizes may be used for X and α without affecting exactness, since each leapfrog update is still a volume-preserving shear and the sequence remains reversible. The code allows this through the parameters `dtau_x` and `dtau_a`.

5 Observables

5.1 The Polyakov Loop

The Polyakov loop is the trace of the holonomy around the thermal circle:

$$P = \frac{1}{N} \text{Tr} \mathcal{P} \exp \left(i \int_0^\beta dt A(t) \right) = \frac{1}{N} \sum_{k=1}^N e^{i\alpha_k}, \quad (24)$$

where the second equality holds in the static diagonal gauge. It is the order parameter of the center symmetry, which shifts all the angles by the same constant, $\alpha_k \rightarrow \alpha_k + c$, and multiplies P by e^{ic} . In the confined phase the α_k are spread uniformly around the circle and $\langle |P| \rangle \rightarrow 0$ as $N \rightarrow \infty$, while in the deconfined phase they clump and $\langle |P| \rangle$ is of order one. Since the overall phase of P wanders freely during the simulation, one measures its modulus:

```
pol = abs(sum(exp(ii*alpha)))/dble(nmat)
```

At finite N , $\langle |P| \rangle$ does not vanish even in the confined phase (for Haar-distributed holonomies it is of order $1/N$), so the transition is smeared out, and it becomes sharp only as $N \rightarrow \infty$.

5.2 Energy

The internal energy is $E = -\partial \ln Z / \partial \beta$. For the bosonic model with $m = 0$, it can be expressed purely in terms of the potential. For a Hamiltonian $H = K + V$ with a potential that is homogeneous of degree four in X , the quantum virial theorem gives $2\langle K \rangle = \langle X \cdot \nabla V \rangle = 4\langle V \rangle$ in every energy eigenstate, and hence $E = 3\langle V \rangle$. With $V = -\frac{N}{4} \sum_{i,j} \text{Tr}[X_i, X_j]^2$ (for $\lambda = 1$), this gives

$$\frac{E}{N^2} = -\frac{3}{4N} \left\langle \frac{1}{\beta} \int_0^\beta dt \sum_{i,j} \text{Tr}[X_i(t), X_j(t)]^2 \right\rangle \rightarrow -\frac{3}{4N} \left\langle \frac{1}{L} \sum_{t=1}^L \sum_{i,j} \text{Tr}[X_i(t), X_j(t)]^2 \right\rangle. \quad (25)$$

The same formula can be obtained by rescaling t and X in the path integral so that β appears only in front of the potential [11, 12]. This estimator is much less noisy than one involving the kinetic term, whose lattice version contains a divergent contribution proportional to $1/a$.

The quantity $\mathcal{H}$ monitored by HMC should not be confused with E : the former is a property of the algorithm, while the latter is a physical observable.

5.3 Extent of Space

A simple measure of the size of the bound state of D0-branes is

$$R^2 = \left\langle \frac{1}{NL} \sum_{t=1}^L \sum_{i=1}^D \text{Tr} X_i(t)^2 \right\rangle, \quad (26)$$

which measures the spread of the eigenvalue distributions of the X_i .

5.4 Statistical Errors

Successive HMC configurations are correlated, so the naive standard error underestimates the true uncertainty. A simple remedy, used in the code, is to divide the measurements into bins that are much longer than the autocorrelation time and to compute the standard error of the bin averages. Autocorrelations are largest near the phase transition (critical slowing down) and for R^2 , which is sensitive to the slow motion of the eigenvalues along the lifted flat directions.

6 Numerical Results

We now illustrate the method with the accompanying code, `BFSS_Model.f90`. All runs use $D = 9$ and $\lambda = 1$, and the results are obtained at a single lattice spacing and a single, small N . They are meant to show that the algorithm works, and not as a precision study.

6.1 Checks of the Algorithm

Table 1 shows the checks of Section 4 for $N = 6$, $L = 8$, $T = 1$ on a thermalized configuration. The analytic forces (17)–(18) agree with symmetric finite differences of the action to seven or eight significant digits. Integrating a trajectory forward and then backward returns to the initial configuration up to 10^{-16} . At fixed trajectory length $n_\tau \Delta\tau = 0.2$, halving $\Delta\tau$ reduces $\Delta\mathcal{H}$ by a factor approaching 4, as expected for a second-order integrator.

$\Delta\tau$	n_τ	$\Delta\mathcal{H}$	ratio
0.02	10	2.17×10^{-1}	–
0.01	20	4.13×10^{-2}	5.2
0.005	40	9.51×10^{-3}	4.3
0.0025	80	2.33×10^{-3}	4.1

Table 1: Energy violation of a single leapfrog trajectory of fixed length, from the same initial configuration and momenta ($N = 6$, $L = 8$, $T = 1$). The last column is the ratio to the previous row.

6.2 Temperature Scan

We simulated $N = 8$ and $L = 12$ at ten temperatures between $T = 0.6$ and $T = 1.5$, with $n_\tau = 40$ and $\Delta\tau = 0.01$ for both X and α . For each temperature we used 500 thermalization trajectories followed by 3000 measured ones, starting from $X = 0$ and uniformly spread angles. The acceptance rate ranges from 0.79 to 0.94, and $\langle e^{-\Delta\mathcal{H}} \rangle$ is compatible with 1 at the percent level at all temperatures, as it should be. Errors are estimated from 20 bins. The scan is run with `run_scan.sh` and analysed with `plot_scan.py`.




Figure 1: Polyakov loop, energy, and extent of space in the bosonic BFSS model for $N = 8$, $L = 12$. The shaded band marks the critical region $0.876 \leq T \leq 0.905$ found in Ref. [11], and the dashed lines show the confined-phase values $E/N^2 = 6.695$ and $R^2 = 2.291$ obtained there at $N = 32$.

The results are shown in Figure 1. The Polyakov loop is small at low temperature and rises steeply between $T = 0.85$ and $T = 0.95$, approaching 1 at high temperature. This is the confinement/deconfinement transition, which appears as a smooth crossover at $N = 8$ and occurs at the location found in the literature [11, 12]. The residual value $\langle |P| \rangle \approx 0.14$ at $T = 0.6$ is a finite- N effect. The energy and the extent of space behave in the same way. In the confined phase they depend only weakly on the temperature, with $E/N^2 \approx 6.50$ – 6.61 and $R^2 \approx 2.25$ – 2.27 for $T \leq 0.8$, close to the $N = 32$ values of Ref. [11]. The remaining differences of about 3% and 2% are consistent with finite- N and finite-lattice-spacing corrections. Above the transition, both quantities grow with the temperature. At high temperature the model becomes weakly coupled and is dominated by the static modes of X_i and A . It then reduces to a ten-matrix Yang–Mills integral, and a scaling argument in that integral gives $E/N^2 \rightarrow 6T$ to leading order. The data approach this slope from above ($E/N^2 \approx 6.9T$ at $T = 1.5$). A careful study would repeat the scan for several values of N and L and extrapolate to $N \rightarrow \infty$ and $a \rightarrow 0$.

7 Discussion and Outlook

The BFSS model provides a non-perturbative window into M-theory and quantum gravity. In these notes we have reviewed its theoretical background and described, step by step, how its bosonic sector is simulated at finite temperature with Hybrid Monte Carlo. The key ingredients are:

- the **static diagonal gauge**, which reduces the gauge field to N angles at the cost of a Faddeev–Popov term and a constraint on the angles;
- the **leapfrog integrator**, whose reversibility and volume preservation make the algorithm exact once combined with the Metropolis test, and whose $O(\Delta\tau^2)$ energy violation controls the acceptance rate;

- **gauge-invariant observables**, such as the Polyakov loop, the energy computed from the commutator term, and the extent of the eigenvalue distribution.

Even at modest N , the simulation exhibits the confinement/deconfinement transition of the bosonic model.

The natural next step is the inclusion of fermions. Integrating out the sixteen-component fermions produces a Pfaffian, which is treated with the rational HMC algorithm. The Pfaffian is in general complex, although its phase was found to have little effect in the regimes studied so far [3,4,13]. With fermions included, the flat directions survive quantum mechanically, and the simulations must be stabilized, for example by restricting the extent of the eigenvalues, a procedure that is justified at large N where the metastable bound state becomes long-lived. Simulations of this kind have reproduced the supergravity prediction (9) and its α' corrections [5,6], providing strong evidence for gauge/gravity duality in a setting where both sides can be computed. Other directions include taking the continuum ($a \rightarrow 0$) and large- N limits systematically, the study of real-time dynamics and chaos in the classical and quantum matrix model, and the use of higher-order symplectic integrators or Fourier acceleration to reduce the cost of the simulations.

References

- [1] T. Banks, W. Fischler, S. H. Shenker, and L. Susskind, “M theory as a matrix model: A conjecture,” *Phys. Rev. D* **55**, 5112 (1997), [arXiv:hep-th/9610043](#).
- [2] N. Itzhaki, J. M. Maldacena, J. Sonnenschein, and S. Yankielowicz, “Supergravity and the large N limit of theories with sixteen supercharges,” *Phys. Rev. D* **58**, 046004 (1998), [arXiv:hep-th/9802042](#).
- [3] K. N. Anagnostopoulos, M. Hanada, J. Nishimura, and S. Takeuchi, “Monte Carlo studies of supersymmetric matrix quantum mechanics with sixteen supercharges at finite temperature,” *Phys. Rev. Lett.* **100**, 021601 (2008), [arXiv:0707.4454](#).
- [4] S. Catterall and T. Wiseman, “Black hole thermodynamics from simulations of lattice Yang-Mills theory,” *Phys. Rev. D* **78**, 041502 (2008), [arXiv:0803.4273](#).
- [5] M. Hanada, Y. Hyakutake, G. Ishiki, and J. Nishimura, “Holographic description of a quantum black hole on a computer,” *Science* **344**, 882 (2014), [arXiv:1311.5607](#).
- [6] E. Berkowitz, E. Rinaldi, M. Hanada, G. Ishiki, S. Shimasaki, and P. Vranas, “Precision lattice test of the gauge/gravity duality at large-N,” *Phys. Rev. D* **94**, 094501 (2016), [arXiv:1606.04951](#).
- [7] E. Witten, “Bound states of strings and p-branes,” *Nucl. Phys. B* **460**, 335 (1996), [arXiv:hep-th/9510135](#).
- [8] B. de Wit, J. Hoppe, and H. Nicolai, “On the quantum mechanics of supermembranes,” *Nucl. Phys. B* **305**, 545 (1988).

- [9] L. Susskind, “Another conjecture about M(atrix) theory,” [arXiv:hep-th/9704080](#).
- [10] W. Taylor, “M(atrix) theory: Matrix quantum mechanics as a fundamental theory,” *Rev. Mod. Phys.* **73**, 419 (2001), [arXiv:hep-th/0101126](#).
- [11] N. Kawahara, J. Nishimura, and S. Takeuchi, “Phase structure of matrix quantum mechanics at finite temperature,” *JHEP* **10**, 097 (2007), [arXiv:0706.3517](#).
- [12] G. Bergner, N. Bodendorfer, M. Hanada, E. Rinaldi, A. Schäfer, and P. Vranas, “Thermal phase transition in Yang-Mills matrix model,” *JHEP* **01**, 053 (2020), [arXiv:1909.04592](#).
- [13] V. G. Filev and D. O’Connor, “The BFSS model on the lattice,” *JHEP* **05**, 167 (2016), [arXiv:1506.01366](#).
- [14] S. Duane, A. D. Kennedy, B. J. Pendleton, and D. Roweth, “Hybrid Monte Carlo,” *Phys. Lett. B* **195**, 216 (1987).
- [15] R. M. Neal, “MCMC using Hamiltonian dynamics,” in *Handbook of Markov Chain Monte Carlo*, Chapman & Hall/CRC (2011).
- [16] N. Metropolis, A. W. Rosenbluth, M. N. Rosenbluth, A. H. Teller, and E. Teller, “Equation of state calculations by fast computing machines,” *J. Chem. Phys.* **21**, 1087 (1953).
- [17] G. E. P. Box and M. E. Muller, “A note on the generation of random normal deviates,” *Ann. Math. Statist.* **29**, 610 (1958).
- [18] W. H. Press, S. A. Teukolsky, W. T. Vetterling, and B. P. Flannery, *Numerical Recipes in Fortran 77: The Art of Scientific Computing*, 2nd ed., Cambridge University Press (1992), Sec. 7.1.
- [19] E. Hairer, C. Lubich, and G. Wanner, *Geometric Numerical Integration: Structure-Preserving Algorithms for Ordinary Differential Equations*, 2nd ed., Springer (2006).